

Exploiting Symmetry in High-Dimensional Dynamic Programming

Mahdi Ebrahimi Kahou, Jesús Fernández-Villaverde, Jesse Perla and Arnav Sood

Boston College

April 8, 2024

- 1 Introduction
- 2 Permutation Invariance
- 3 Concentration of Measure
- 4 Model
- 5 Deep Learning
- 6 Extensions
- 7 Discussion

Introduction

- ◀ A new method for solving high-dimensional dynamic programming problems with a finite but large number of heterogeneous agents. Accurate and quick.
- ◀ Exploiting symmetry.
- ◀ Concentration of measure.
- ◀ Design and train deep learning architectures that exploit symmetry and concentration of measure.
- ◀ Use multi-firm investment under uncertainty model to demonstrate the performance of the method.

Permutation Invariance

- ◀ Exploit symmetry in the solution of dynamic programming problems to reduce dimensionality.
- ◀ Permutation-invariant function:

$$f(x, \pi X) = f(x, X)$$

π is a N -dimensional permutation matrix

- ◀ Proposition 2: Representation of permutation-invariant functions

$$f(x, X) = \rho(x, \frac{1}{N} \sum_{i=1}^N \phi(X_i))$$

$$\rho : R^{L+1} \rightarrow R, \phi : R \rightarrow R^L$$

- ◀ How to exploit it?

Concentration of Measure

- Can calculate high-dimensional expectations using single Monte Carlo draw.
- Proposition 3: Concentration of measure when expected gradients are bounded in N

$$P(|f(z^1) - E[f(z)]| \geq \epsilon) \leq \frac{\rho(\Sigma)}{\epsilon^2} \frac{1}{N}$$

Provide an upper bound on the error of the evaluation of the expectation.
It is usually better to approximate $E[f(z)]$ by $f(z^1)$ than by $f(\mathbf{0}_N)$

- How to exploit it?

Model

Investment under uncertainty with many firms:

$$v(x, X) = \max_u p(X)x - \frac{\gamma}{2}u^2 + \beta E[v(x', X')]$$

$$s.t. x' = (1 - \delta)x + u + \sigma w + \eta \omega$$

$$X'_i = (1 - \delta)X_i + \hat{u}(X_i, X) + \sigma W_i + \eta \omega, \text{ for } i \in \{2, \dots, N\}$$

$$X'_1 = (1 - \delta)X_1 + \hat{u}(X_1, X) + \sigma w + \eta \omega$$

where $p(X)$ (the inverse demand function) is assumed to be:

$$p(X) = \alpha_0 - \alpha_1 \frac{1}{N} \sum_{i=1}^N x_i^\nu$$

Deep Learning

- ◀ Approximate ρ , ϕ and L using a deep learning architecture $F(\theta)$.

$$u(x, X) = \rho(x, \frac{1}{N} \sum_{i=1}^N \phi(X_i))$$

Baseline case: $\phi(\text{Identity})$, $\phi(\text{Moments})$ and $\phi(\text{ReLU})$ with two layers each with 128 nodes. 49.2K, 49.8K, and 66.8K parameters respectively, regardless of N.

- ◀ Train the network by minimizing the Euler residuals:

$$\varepsilon(x, X; \theta) \equiv \gamma u(x, X; \theta) - \beta E[P(X') + \gamma(1 - \delta)u(x', X'; \theta)]$$

Deep Learning

Case I: $\nu = 1$

Analytical solution: $u(X) = H_0 + \frac{1}{N} H_1 \sum_{i=1}^N x_i$




Figure 3: Comparison between the LQ-regulator solution and our three deep learning architectures for the case with $\nu = 1$ and $N = 128$.

Deep Learning

Accuracy of the approximation:




Figure 4: The Euler residuals for $\nu = 1$ and $N = 128$ for $\phi(\text{Identity})$, $\phi(\text{Moments})$, and $\phi(\text{ReLU})$. The dark blue curve shows the average residuals along equilibrium paths for 256 different trajectories. The shaded areas depict the 2.5th and 97.5th percentiles.

The ReLU architecture is especially stable. The extra parameters of this architecture help the function to generalize very consistently.

Deep Learning

Is it practical to implement?




Figure 5: Performance of the $\phi(\text{ReLU})$ for different N .

Computation time: of order $O(1)$ for reasonable N .

Deep Learning

Explore with different network architectures:

Table 1: Performance of different networks in solving case I: $\nu = 1$

group	description	Time (s)	Params (K)	Train MSE (ε)	Test MSE (ε)	Val MSE (ε)	Policy Error ($ u - u_{\text{ref}} $)	Policy Error ($\frac{ u - u_{\text{ref}} }{u_{\text{ref}}}$)
$\phi(\text{Identity})$	Baseline	42	49.4	4.1e-06	3.3e-07	3.3e-07	2.9e-05	0.10%
	Thin (64 nodes)	33	12.4	3.7e-06	2.7e-07	2.7e-07	3.4e-05	0.10%
$\phi(\text{Moments})$	Baseline	55	49.8	1.4e-06	7.6e-07	7.6e-07	2.8e-05	0.09%
	Moments (1,2)	211	49.5	2.4e-06	1.1e-06	2.3e-06	4.4e-05	0.14%
	Very Shallow(1 layer)	241	0.6	1.1e-05	8.4e-06	7.9e-06	1.1e-02	34.00%
	Thin (64 nodes)	82	12.6	1.6e-06	9.1e-07	9.2e-07	3.8e-05	0.12%
$\phi(\text{ReLU})$	Baseline	107	66.8	3.7e-06	3.3e-07	3.3e-07	2.7e-05	0.09%
	L = 2	86	66.3	1.3e-05	2.1e-07	2.2e-07	2.6e-05	0.08%
	L = 16	91	69.9	5.5e-06	1.5e-07	1.5e-07	2.1e-05	0.07%
	Shallow(ϕ : 1 layer, ρ : 2 layers)	79	17.7	2.0e-06	5.5e-07	5.5e-07	3.2e-05	0.11%
	Deep(ϕ : 4 layers, ρ : 8 layers)	242	165.1	2.1e-03	2.2e-03	2.1e-03	2.7e-03	8.50%
	Thin(ϕ , ρ : 64 nodes)	87	17.0	1.1e-05	4.5e-07	4.5e-07	3.0e-05	0.10%

- Very shallow in $\phi(\text{Moments})$ finds a local minimum. Deep in $\phi(\text{ReLU})$ is unsatisfactory. Explore different architectures.
- Results are good when we use a higher dimension of L .

Deep Learning

Case II: $\nu > 1$

No analytical solution.




Figure 6: The Euler residuals for $\nu = 1.5$ and $N = 128$ for $\phi(\text{Moments})$ and $\phi(\text{ReLU})$. The dark blue curve shows the average residuals along equilibrium paths for 256 different trajectories. The shaded areas depict the 2.5th and 97.5th percentiles.

Extensions

The tools are useful for solving any high-dimensional functional equations with some degree of symmetry, especially when these equations contain high-dimensional expectations.

- ◀ Decreasing returns to scale
- ◀ Multiple productivity types
- ◀ Complex idiosyncratic states
- ◀ Global solutions with transitions and aggregate shocks

Discussion

- ◀ Solve high-dimensional dynamic programming problems in minutes.
- ◀ Double-descent: by increase the number of parameters, we can both fit the data perfectly and achieve outstanding generalization.
- ◀ Model selection since results are sensitive to different network architectures.